

Project Explanation

Overview

The **Insurance Agent** project is a full-stack AI-driven application that provides auto and home insurance quotes through a conversational interface. It combines a FastAPI backend, a React/Vite frontend, and a LangGraph-based workflow (the "brain") that orchestrates LLM reasoning, tool calls, and RAG search.

Repository Structure

```
insurance_agent/
  |- backend/           # FastAPI server and AI logic
    |   main.py          # API endpoints, session handling
    |   memory.py        # Simple in-memory store for session data
    |   provider.py      # Gemini provider abstraction
    |   langgraph_agent.py # LangGraph workflow definition
    |   rag_system.py    # Chroma vector store + Gemini embeddings
    |   document_analyzer.py # Gemini Vision document analysis
    |   requirements.txt
    |   test_*.py         # Unit & integration tests
    |   __init__.py
  |- frontend/          # React UI built with Vite
    |   src/
    |     |   components/ChatInterface.jsx # Chat UI component
    |     |   App.jsx
    |     |   main.jsx
    |   index.html
    |   index.css          # Custom palette, Google Font "Inter"
    |   package.json
    |   vite.config.js
  .gemini/brain/        # Artifacts generated during development
    |   implementation_plan.md
    |   task.md
    |   walkthrough.md
    |   agent_architecture.md
```

Key Technologies

- **FastAPI + Uvicorn** – backend web framework and server.
- **React + Vite** – modern, fast frontend development.
- **Google Gemini 1.5 Flash** ([gemini-2.0-flash-exp](#)) – LLM for reasoning and vision.
- **LangGraph** – visual state-graph workflow (nodes, edges, conditional routing).
- **Chroma DB + Gemini embeddings** – vector store for Retrieval-Augmented Generation (RAG).
- **LangChain** – tool orchestration and message handling.

- **dotenv** – loads the `GEMINI_API_KEY` from `.env`.
- **Python-multipart** – handles file uploads for document analysis.
- **Testing** – `unittest` + FastAPI `TestClient` for API and LangGraph tests.

How It Works

1. **User** types a message in the React UI → POST `/api/chat`.
2. **FastAPI** stores the message in an in-memory session dict.
3. **(Future)** `main.py` forwards the session state to `agent_graph.invoke(state)`.
 - The graph runs `gather_info`, decides whether to `search_knowledge`, `calculate_quote`, or loop for more info.
 - Tool nodes call the premium-calculation functions.
4. If a knowledge query is detected, **RAG** (`rag_system.search_knowledge`) retrieves relevant docs from Chroma and injects them into the prompt.
5. **Gemini** generates the final answer, which is returned to the UI.
6. The UI displays the response; screenshots of the interaction are saved in the artifact folder for the walkthrough.

Testing & Verification

- Unit tests for memory, provider, and API endpoints.
- `test_langgraph_brain.py` validates the LangGraph state transitions.
- End-to-end UI tests (via browser sub-agent) captured screenshots for refusal and auto-insurance flows.

This file provides a concise, printable summary of the project for documentation or workshop hand-outs.