

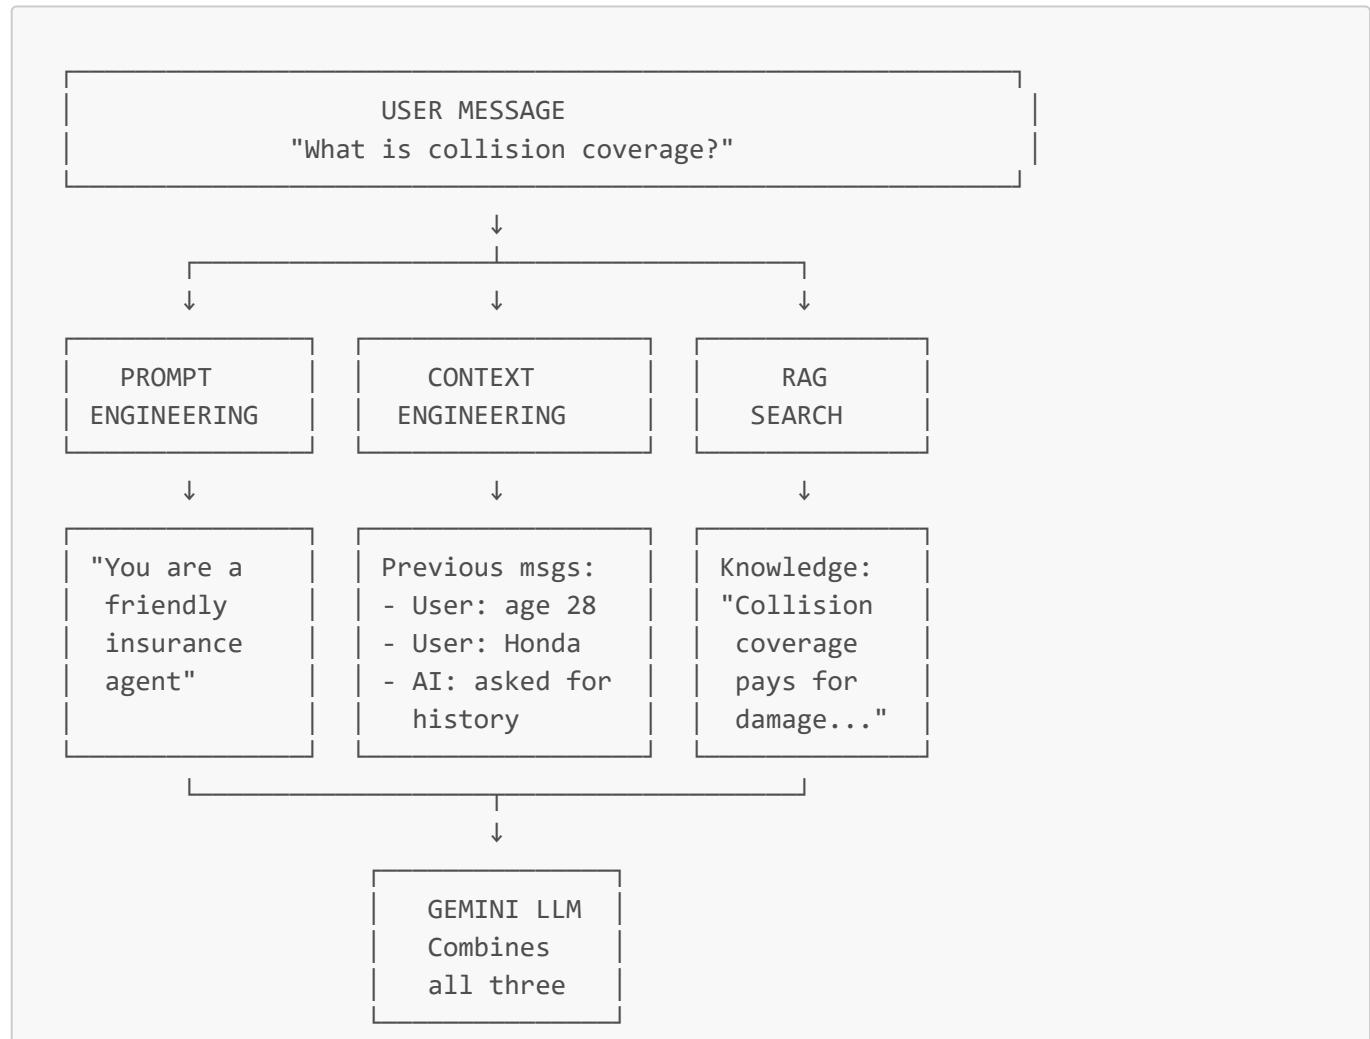
# Context Engineering vs Prompt Engineering vs RAG

## Understanding the Three Pillars of AI Intelligence

### Quick Comparison

Aspect	Prompt Engineering	Context Engineering	RAG (Retrieval Augmented Generation)
<b>What</b>	Design AI personality & behavior	Provide relevant info for THIS conversation	Search knowledge base for facts
<b>When</b>	Once, at design time	Every message	When user asks questions
<b>Who</b>	Orchestrators design	System provides automatically	System searches automatically
<b>Where</b>	System prompt	Conversation history	External knowledge base
<b>Example</b>	"You are a friendly agent"	"User said age is 28"	"Collision coverage definition"

### Visual Comparison



↓

RESPONSE

"Collision  
coverage..."

## 1 PROMPT ENGINEERING

### What It Is

Designing the AI's **personality, role, and behavior** that stays constant across all conversations.

### Analogy

Like hiring an employee and giving them a job description.

### In the Workshop

#### Part 2 (15 min): Orchestrators design the agent's personality

##### Example Prompt:

You are "Alex", an expert insurance agent powered by AI.

Your personality:

- Friendly and professional
- Patient and helpful
- Explains complex terms simply

Your role:

1. Help customers get accurate quotes
2. Answer questions about coverage
3. Gather required information

Guidelines:

- Ask 1-2 questions at a time
- Explain clearly when asked
- Calculate quote when ready

### Code Location

backend/system\_prompt.py

```
INSURANCE_AGENT_PROMPT = """
You are "Alex", an expert insurance agent...
"""
```

## When It's Used

- **Once** when agent starts
- **Same for all users**
- **Defines WHO the agent is**

## Orchestrator Activity

1. Open AI Studio
  2. Design personality
  3. Test with conversations
  4. Refine based on results
  5. Share with implementers
- 

## 2 CONTEXT ENGINEERING

### What It Is

Providing the AI with **relevant information from THIS specific conversation** so it remembers what was said.

### Analogy

Like taking notes during a meeting so you remember what was discussed.

### In the Workshop

#### Part 3 (20 min): Implementers build conversation memory

##### Example Context:

```
Conversation so far:  
User: "I need car insurance"  
AI: "I'd be happy to help! What's your age?"  
User: "I'm 28"  
AI: "Great! What vehicle do you drive?"  
User: "2020 Honda Civic"  
AI: "Perfect! How many years have you been licensed?"
```

### Code Location

backend/main.py (lines 132-142)

```
# Get or create session  
if session_id not in sessions:  
    sessions[session_id] = {  
        "messages": [], # ← Context: conversation history  
        "user_info": {}, # ← Context: extracted data  
        "insurance_type": None, # ← Context: what they need
```

```
        "quote_result": None,  
        "knowledge_context": "",  
        "next_action": "gather_info"  
    }  
  
session = sessions[session_id]  
  
# Add user message to history  
session["messages"].append(HumanMessage(content=request.message))
```

## When It's Used

- Every message in the conversation
- Unique per user
- Defines WHAT was said

## Implementer Activity

1. Store conversation history
2. Track extracted information
3. Pass to LLM with each message
4. Agent remembers context

## Why It Matters

### Without context:

```
User: "I'm 28"  
AI: "What's your age?" ← Forgot user just said it!
```

### With context:

```
User: "I'm 28"  
AI: "Great! What vehicle do you drive?" ← Remembers age
```

## ③ RAG (Retrieval Augmented Generation)

### What It Is

Searching a knowledge base for facts and information to answer user questions accurately.

### Analogy

Like looking up information in a company handbook before answering.

### In the Workshop

## Part 4 (20 min): Orchestrators write FAQs, Implementers build search

### Example Knowledge Base:

Q: What is collision coverage?

A: Collision coverage pays for damage to YOUR vehicle when you hit another vehicle or object, regardless of who's at fault.

Q: What is comprehensive coverage?

A: Comprehensive coverage pays for damage to your vehicle from non-collision events like theft, vandalism, weather, or fire.

Q: How can I lower my premium?

A: Common ways include: bundling policies (10-25% off), maintaining a clean driving record (15-30% off), and increasing deductible.

## Code Location

[backend/rag\\_system.py](#)

```
# Knowledge base (from orchestrators!)
INSURANCE KNOWLEDGE = {
    "auto_coverage": [
        "Collision coverage pays for damage to YOUR vehicle...",
        "Comprehensive coverage pays for damage from theft...",
        "Liability coverage pays for damage YOU cause..."
    ],
    "discounts": [
        "Bundling policies saves 10-25%...",
        "Clean driving record saves 15-30%"
    ]
}

def search_knowledge(query: str, k: int = 2):
    """Search knowledge base for relevant information"""
    vectorstore = Chroma(...)
    results = vectorstore.similarity_search(query, k=k)
    return results
```

## When It's Used

- **When user asks questions** ("What is X?")
- **Searches external knowledge**
- **Provides FACTS to the agent**

## Orchestrator Activity

1. Write 10 insurance FAQs

2. Cover: coverage types, discounts, general
3. Keep answers clear
4. Share with implementers

## Implementer Activity

1. Add FAQs to **INSURANCE KNOWLEDGE**
2. Initialize vector database
3. Implement search function
4. Integrate with chat endpoint

## How It Works

```
User asks: "What is collision coverage?"  
↓  
System detects question keyword ("what is")  
↓  
RAG searches knowledge base  
↓  
Finds: "Collision coverage pays for damage..."  
↓  
Adds to context for LLM  
↓  
LLM uses fact to answer accurately
```

## 🔗 How They Work Together

### Example Conversation

**User:** "Hi, I need insurance"

**What happens:**

1. **Prompt Engineering:** Agent uses personality ("friendly, helpful")
2. **Context Engineering:** Stores message in history
3. **RAG:** Not triggered (no question asked)

**Response:** "I'd be happy to help! Are you looking for auto or home insurance?"

---

**User:** "Auto insurance. What's the difference between collision and comprehensive?"

**What happens:**

1. **Prompt Engineering:** Agent uses guideline ("explain clearly")
2. **Context Engineering:** Remembers user wants auto insurance
3. **RAG:** Searches knowledge base for "collision" and "comprehensive"

**RAG finds:**

- Collision: pays for damage to YOUR vehicle when you hit something
- Comprehensive: pays for damage from theft, weather, etc.

**Response:** "Great question! Collision coverage pays for damage to YOUR vehicle when you hit another vehicle or object. Comprehensive coverage pays for damage from non-collision events like theft, vandalism, or weather. Both are optional but recommended!"

---

**User:** "I'm 28, drive a 2020 Honda Civic"

**What happens:**

1. **Prompt Engineering:** Agent uses guideline ("ask 1-2 questions at a time")
2. **Context Engineering:** Stores age=28, vehicle=2020 Honda Civic
3. **RAG:** Not triggered (no question)

**Response:** "Perfect! How many years have you been licensed, and have you had any accidents or violations in the last 3 years?"

---

## Workshop Timeline

### Part 2: Prompt Engineering (15 min)

**Orchestrators:** Design agent personality in AI Studio **Implementers:** Code `system_prompt.py` **Together:** Test personality

### Part 3: Context Engineering (20 min)

**Orchestrators:** Create test cases **Implementers:** Build LangGraph with state management **Together:** Verify agent remembers context

### Part 4: RAG (20 min)

**Orchestrators:** Write 10 FAQs **Implementers:** Implement vector search **Together:** Test knowledge retrieval

---

## Key Differences

### Prompt Engineering

- **Static:** Same for all users
- **Design time:** Created once
- **Defines:** WHO the agent is
- **Example:** "You are friendly and professional"

### Context Engineering

- **Dynamic:** Unique per conversation
- **Runtime:** Updates every message

- **Defines:** WHAT was said
- **Example:** "User said age is 28"

## RAG

- **On-demand:** Only when needed
  - **Query time:** Searches when user asks
  - **Defines:** FACTS from knowledge base
  - **Example:** "Collision coverage definition"
- 

## 💡 When to Use Each

### Use Prompt Engineering When:

- Defining agent personality
- Setting conversation guidelines
- Specifying required information
- Establishing tone and style

### Use Context Engineering When:

- Remembering user information
- Tracking conversation flow
- Avoiding repeated questions
- Maintaining state across messages

### Use RAG When:

- User asks factual questions
  - Need accurate, up-to-date information
  - Have large knowledge base
  - Want to avoid hallucinations
- 

## ✍️ Hands-On Exercise

### Test All Three

**Scenario:** User asks about insurance

**Your Turn:**

1. **Prompt Engineering:** Design personality

You are [NAME], a [ROLE] with [PERSONALITY].

2. **Context Engineering:** Track this conversation

```
User: "I need car insurance"
AI: "What's your age?"
User: "28"
AI: [Should remember age, ask next question]
```

### 3. RAG: Write a FAQ

```
Q: What is liability coverage?
A: [Your answer]
```

## Code Examples

### 1. Prompt Engineering

```
# backend/system_prompt.py
INSURANCE_AGENT_PROMPT = """
You are "Alex", an expert insurance agent.
Your personality: Friendly, professional, helpful
"""
```

### 2. Context Engineering

```
# backend/main.py
session = {
    "messages": [
        HumanMessage("I need insurance"),
        AIMessage("What's your age?"),
        HumanMessage("I'm 28") # ← Context stored
    ],
    "user_info": {"age": 28} # ← Extracted context
}
```

### 3. RAG

```
# backend/rag_system.py
INSURANCE KNOWLEDGE = {
    "auto_coverage": [
        "Collision coverage pays for damage to YOUR vehicle..."
    ]
}

# Search when user asks
```

```
results = search_knowledge("What is collision coverage?")
# Returns: "Collision coverage pays for damage..."
```

## Summary Table

Feature	Prompt Eng	Context Eng	RAG
<b>Frequency</b>	Once	Every message	On-demand
<b>Scope</b>	All users	Per user	Per question
<b>Storage</b>	Code	Memory	Database
<b>Created by</b>	Orchestrators	System	Orchestrators + System
<b>Purpose</b>	Define behavior	Remember conversation	Provide facts
<b>Example</b>	"Be friendly"	"User is 28"	"Collision = ..."

## Key Takeaways

1. **Prompt Engineering** = WHO the agent is (personality)
2. **Context Engineering** = WHAT was said (memory)
3. **RAG** = FACTS from knowledge base (search)
4. **All three work together** to create intelligent conversations
5. **Orchestrators design** prompts and knowledge
6. **Implementers build** context management and RAG
7. **Together** they create a smart agent

## Workshop Resources

- **Prompt Engineering:** [docs/PROMPT\\_ENGINEERING\\_GUIDE.md](#)
- **Context Engineering:** See LangGraph agent state
- **RAG:** [backend/rag\\_system.py](#)

**Remember:** These aren't competing approaches - they complement each other! 